

Introduction to machine learning

Maksim KretoV

Lecture 7: Training deep neural networks

5vision, 2017

Course information I

Course

10 lectures + 2 seminars; February-May 2017.

Schedule and up-to-date syllabus

<https://goo.gl/xExEuL>

Contact information and discussion

Maksim Kretov (kretovmk@gmail.com)

Alexey Seleznev (a.o.seleznev@gmail.com)

Slack group: <https://miptmlcourse.slack.com>

to get an invite, send e-mail to kretovmk@gmail.com.

Course information II

Grading policy

3 practical assignments, each gives 25, exam: 25

=> Total: 100. Pass: ≥ 50

Maksim Kreto

Lectures 1-7, PA1 and exam

Alexey Seleznev

Lectures 8-12, PA2-3 and exam

Average 18 students,
70% recurrent



4 PA1 done



?

Plan of the course

Math and basics of ML	(1-2)	}	<i>Theoretical tasks</i>		
Some of ML methods	(3)				
Seminar on ML basics	(4)				
Basics of neural networks	(5)	}	<i>+Practical tasks</i>		
Deep learning overview	(6)				
<u>Training deep networks</u>	(7)			← Today	
DL for Computer Vision	(8-9)			}	Solving more complex ML tasks using NNs
DL for time series prediction	(10-11)				
<i>Concluding seminar</i>	(12)				

Plan for the lecture

A. Previous lecture

1. ERM framework
2. Deep learning

B. Improving convergence of BP

1. Cross-entropy
2. Weights initialization

C. Gradient descent

1. Stochastic approximation
2. Advanced GD: momentum, adagrad etc.

D. Regularization: Dropout, Batch normalization.

E. Baseline models

F. Practical assignments

A.1 Previous lectures: ERM framework

Empirical risk minimization approach (ERM)

Formula for fitting the model within ERM framework:

$$\theta^{opt} = \operatorname{argmin}_{\theta} \frac{1}{N} \sum_{n=1}^N L(y_n, \hat{y}_n) + \lambda \Omega(\theta)$$

$L(y_n, f(\mathbf{x}_n, \theta))$ is loss function; $\hat{y}_n = f(\mathbf{x}_n, \theta)$ is prediction

$\Omega(\theta)$ is a regularizer => **learning converted into optimization task!**

Training neural networks with ERM. Probabilistic interpretation:

Maximizing likelihood of correct class in predicted distribution.

A.2 Previous lectures: Deep learning

Just an informal definition of “deep” networks

Networks with up to 3 (2 hidden) layers → shallow

More than 3 layers → deep

Traditional methods: local smoothness assumption

Deep learning methods: complement with “compositionality” prior.

Deep Learning

Machine learning algorithms based on learning multiple levels of representation / abstraction.*

B. Improving convergence

Why do we need special training protocol for deep networks?

- Loss function is complex and non-convex (underfitting)
- Lots of parameters, more than examples in training set (overfitting)
- Huge datasets (slow, don't fit in memory)
- NN-specific problems (vanishing gradient)

B.1 Improving convergence: Cross-entropy

Cross-entropy cost

$$L(Y, f(\mathbf{X}, \theta)) = -\frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K y_{nk} \ln \hat{y}_{nk}$$

y_{nk} are true labels and \hat{y}_{nk} are predictions.

N – number of training examples; K is number of classes in classification task.

Each term is log probability of sampling true label from predicted distribution.

Such loss function meets our criteria:

- connected to a clear proxy such as accuracy
- smooth and easy to differentiate (accuracy is not smooth!)
- probabilistic interpretation: MLE estimation $f(\mathbf{x}, \theta) = p(y|\mathbf{x}, \theta)$ or statistics


B.1 Improving convergence: Cross-entropy

Cross-entropy + sigmoid output

Cross-entropy allows to avoid vanishing gradient for some activation functions in neural networks.

$$\frac{\partial L_{CE}(y_n, \hat{y}_n)}{\partial \omega_i} = x_{ni}(\hat{y}_n - y_n)$$

$$\frac{\partial L_{MSE}(y_n, \hat{y}_n)}{\partial \omega_i} = x_{ni}(\hat{y}_n - y_n)\sigma'(z_n)$$



If weighted input to neuron is close to 1, then $\sigma'(z_n)$ is close to 0 and learning starts to slowing down (“saturation”).

B.1 Improving convergence: Cross-entropy

Cross-entropy + softmax output

Also provides non-saturated architecture:

$$\log \text{softmax}(z)_i = z_i - \log \sum_j \exp z_j \quad \textit{Exercise: check it and MSE+softmax.}$$

Other considerations

1. In practice, logits are used as outputs: $z_i = \omega_i a_{i-1} + b_i$ (weighted input)

2. Multi-label classification:

Cross entropy can be recorded separately for each neuron with sigmoid:

$$L(Y, f(\mathbf{X}, \theta)) = -\frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K (y_{nk} \ln \hat{y}_{nk} + (1 - y_{nk}) \ln(1 - \hat{y}_{nk}))$$

B.2 Improving convergence: Weights init.

Weights initialization

Is the simplest solution possible?

If initialize with zeroes => equal gradient for all neurons.

Simple random initialization (“breaking symmetry”):

Use independent Gaussian RV with $\mu = 0, \sigma = 1$ (or uniform noise).

Problems with that approach:

Standard deviation increases as $\sim \sqrt{n}$, where n is number of inputs
=> may cause saturation of neurons or gradient exploding

Solution:

Initialize with variation $\sim 1/\sqrt{n_{inputs}}$.

B.2 Improving convergence: Weights init.

Weights initialization examples

1. Uniform distributions

$$\omega_{ij} \sim U\left(-\frac{1}{\sqrt{m}}, -\frac{1}{\sqrt{m}}\right)$$

$$\omega_{ij} \sim U\left(-\frac{\sqrt{6}}{\sqrt{m+n}}, -\frac{\sqrt{6}}{\sqrt{m+n}}\right)$$

“Glorot uniform”

2. Normal distribution

$$\omega_{ij} \sim N\left(0, \frac{2}{m}\right) \quad \text{“He normal”}$$

3. Initialization with random orthogonal matrices

5 minute break..

Questions?

C.1 Gradient descent: Stochastic appr.

Robbins-Monro algorithm

Task is to compute root of $f(\theta) = E_{\tau \sim p(\tau)}[\varphi(\theta, \tau)]$

If $p(\tau)$ is not known \Rightarrow algorithm:

$\theta_n = \theta_{n-1} - \mu_n \varphi(\theta_{n-1}, \tau_n)$ converges to root of $f(\theta)$ under conditions:

$$\sum_n \mu_n^2 < \infty \text{ and } \sum_n \mu_n \rightarrow \infty$$

Adaptation to ML tasks

$$\tilde{L}(\theta) = E_{x,y \sim p_{data}}[L(\theta, x, y)]$$

We interested in extremum (i.e. root of the corresponding gradient):

$$\nabla \tilde{L}(\theta) = E_{x,y \sim p_{data}}[\nabla L(\theta, x, y)] \quad \Rightarrow \quad \theta_n = \theta_{n-1} - \mu_n \nabla L(\theta_{n-1}, x_n, y_n)$$

C.2 Gradient descent: Advanced SGD

Gradient Descent flavors:

Batch GD: $\theta^{(i)} = \theta^{(i-1)} - \alpha_i \nabla_{\theta} L(Y, f(\mathbf{X}, \theta))$ [all training set]

Redundant computations (many correlated examples in dataset)

Online GD: $\theta^{(i)} = \theta^{(i-1)} - \alpha_i \nabla_{\theta} L(y, f(\mathbf{x}, \theta))$ [one example]

Variance is too big and not speed up from matrix computations

Mini-batch GD: $\theta^{(i)} = \theta^{(i-1)} - \alpha_i \nabla_{\theta} L(Y_m, f(\mathbf{X}_m, \theta))$ [mini-batch]


Most widely-used now

C.2 Gradient descent: Advanced SGD

Selecting better learning schedule



1. Momentum

$$v_t = \gamma v_{t-1} + \alpha \nabla_{\theta} L(\theta)$$
$$\theta = \theta - v_t$$



2. Nesterov accelerated gradient

$$v_t = \gamma v_{t-1} + \alpha \nabla_{\theta} L(\theta - \gamma v_{t-1})$$
$$\theta = \theta - v_t$$



* Images from

<http://sebastianruder.com/optimizing-gradient-descent/>

<http://cs231n.github.io/neural-networks-3/#update>

C.2 Gradient descent: Advanced SGD

Annealing (scheduling) learning rate

1. Step decay:

Reduce the learning rate by some factor every few epochs.

2. Exponential decay

$$\alpha_t = \alpha_0 \exp(-kt)$$

3. 1/t decay

$$\alpha_t = \alpha_0 / (1 + kt)$$

4. Linear decay

Reason: to prevent “bouncing” around minimums.

C.2 Gradient descent: Advanced SGD

Per-parameter adaptive learning rate methods

1. Adagrad

$$c = c + [\nabla_{\theta} L(\theta)]^2 \quad [this is a vector of sum of squared gradients]$$

$$\theta^{(i)} = \theta^{(i-1)} - \alpha_i \nabla_{\theta} L(\theta) / (\sqrt{c} + \varepsilon) \quad [\varepsilon is for numerical stability]$$


2. RMSprop

$$c = \gamma c + (1 - \gamma) [\nabla_{\theta} L(\theta)]^2 \quad [this is a vector of av. squared gradients]$$

$$\theta^{(i)} = \theta^{(i-1)} - \alpha_i \nabla_{\theta} L(\theta) / (\sqrt{c} + \varepsilon) \quad [\varepsilon is for numerical stability]$$

C.2 Gradient descent: Advanced SGD

Different SGD techniques, demonstration



3. Adam

$$b = \alpha c + (1 - \alpha) \nabla_{\theta} L(\theta)$$

$$c = \gamma c + (1 - \gamma) [\nabla_{\theta} L(\theta)]^2$$

$$\theta^{(i)} = \theta^{(i-1)} - \alpha_i b / (\sqrt{c} + \varepsilon)$$

D. Regularization

Overfitting problem

To prevent neural network from “memorizing” inputs.

[Number of parameters >> with numbers of training examples]

Techniques

1. Early stopping: stop training once validation error starts to increase.
2. L2 regularization

L2 regularization

Regularization L2: control over weights in neural network.

L2 regularization: add term $\frac{\lambda}{2n} \|\omega\|^2$ to cost function ($\lambda > 0$).

D. Regularization


L2 regularization (“weight decay”)

Ideas behind regularization:

- to make model “simpler”

- to make model more stable to random noise in input data

- fewer weights \Rightarrow smaller VC dimension




D. Regularization

Dropout:

1. Randomly and temporarily delete neurons
[forward and backward passes through modified NN]
2. Repeat for another mini-batch (select new subset of neurons for deletion).

Intuition behind procedure:

Training different neural networks



D. Regularization

Batch normalization:

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1...m}\}$;
Parameters to be learned: γ, β
Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$
$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Algorithm 1: Batch Normalizing Transform, applied to activation x over a mini-batch.

Normalizing activations of layer.
 γ, β are learnable parameters.

Robust to “bad” initializations of weights.

Address internal covariance shift problem for deeper layers of NN.

D. Regularization

Other regularization techniques

1. Batch normalization
2. L1 regularization (weights selection)
3. Dataset augmentation (object recognition in CV, NLP)
4. Adding noise to inputs / hidden / weights / targets and labels smoothing
5. Pre-training

E. Baseline models

After the simplest approaches tried...

Neural networks

1. Type of input determines general structure of NN:
 - Images => CNN with common sets of layers
 - Sequential data => LSTM / GRU with FC layers
 - Other => few FC layers
2. Activation function: ReLU (leaky ReLU)
3. Optimizer: Adam (RMSprop, SGD with momentum and decaying lr)
4. Regularizer: Batch Normalization and early stopping (Dropout)

F. Practical assignments

Practical assignment #2: starting 12 Apr

Working with simplified version of Diabetic Retinopathy Detection competition on Kaggle.

Work with 512x512 images. Images already preprocessed. Baseline solution is provided.

Ref: <https://github.com/5vision/miptmlcourse>

Practical assignment #3: starting 19 Apr

Detection of types of physical activities.

Next week

Lecture 8 “Convolutional neural networks”

1. Motivation and premises
2. Key features of CNN architecture
3. Examples of CNNs

D. Homework

For all

1. Reading Ch.7-8 in [2].
2. Practical assignment #2.

Some cool reading:

<http://sebastianruder.com/optimizing-gradient-descent/>

<http://distill.pub/2017/momentum/>

<http://blog.smola.org/post/4110255196/real-simple-covariate-shift-correction>

Refs

1. Thorough review of relevant math topics:

<http://info.usherbrooke.ca/hlarochelle/ift725/review.pdf>

2*. Ian Goodfellow, Yoshua Bengio and Aaron Courville, Deep Learning.

3. Kevin P. Murphy, Machine Learning: A probabilistic perspective.

4. David Barber, Bayesian Reasoning and Machine Learning.

5. Sergios Theodoridis, Machine Learning: A Bayesian and optimization perspective.

6*. See also refs in practical assignment and online courses, especially one from Hugo Larochelle (presentation from lecture 1).